

Beyond Ambiguity: Can NLP Detect Romantic Undertones in The Iliad and The Song of Achilles?

Table of Contents

Introduction	3
Research Objectives	4
Methodology	5
Data and Preprocessing	5
Sentiment Analysis	6
Introduction to transformers models and GoEmotions	6
Experiment design and sentiment analysis process	8
Experiment objective and interpretation	9
Relationship Classification and Topic modelling	10
Principle of Zero-shot Learning and BART-large-MNLI Model	10
Experiment objective and the design with implement process	11
Experiment interpretation and significance	11
Topic Modeling	12
Finding and Conclusion	13
Sentiment Analysis	13
NLP detection of romantic undertones and implications	16
NLP in humanities analysis : strengths and limitations	19
Bibliography	21

Introduction

The nature of Achilles and Patroclus' relationship in *The Iliad* has been debated by scholars for a long time, shifting between interpretations of deep friendship and romantic affection. Scholars like Gregory Nagy (1979) have highlighted the Homeric idea of *philia*, describing their bond as one rooted in heroic camaraderie rather than romantic love. On the other hand, thinkers such as David M. Halperin (1990) caution against applying modern sexual categories to ancient narratives. Since the Greek epic doesn't provide a clear label for their bond, emotional hints often remain wrapped in poetic subtlety. A key term in this debate is "philtatos", a word Achilles uses for Patroclus that roughly translates to "most beloved." Eva Cantarella (1992) has argued that this, along with the extremity of Achilles' grief, could indicate a connection that goes beyond mere friendship. Traditional interpretations like these mostly rely on close readings and classical analysis, often shaped by scholars' subjective perspectives.

By contrast, *The Song of Achilles*, a modern retelling by Madeline Miller presents their relationship directly through a romantic lens, told from Patroclus' first-person perspective. After investigating it, I consider the novel doesn't just reinterpret the myth but intentionally reshapes it as a romantic tragedy, connecting ancient emotion to modern readers' sensibilities. In this project, both texts are analyzed in their English versions, allowing a consistent comparison using a Transformer-based NLP model. For *The Iliad*, Robert Fagles' translation was selected because it keeps much of the original's poetic tone and ambiguity, while *The Song of Achilles* represents a widely accepted modern romantic portrayal, making it suitable for testing how NLP models handle emotional nuance in literature.

While most traditional approaches depend heavily on interpretation, the growth of Natural Language Processing opens up new ways to explore literature through data.

Transformer-based models like BERT and GPT are now used to detect tone, sentiment, and relationship cues in natural language (Underwood, 2019). Research by Bamman, Underwood, and Smith (2020) has shown these models can identify emotions and character relationships in historical texts. But when it comes to literature, especially poetry or classical texts, the NLP might face a real challenge for detecting the literature content. From the previous Underwood (2019)'s research has noted, large language models often struggle with figurative language and subtle emotional layers, especially when those feelings are conveyed through metaphors or culturally specific expressions.

Because of these traits, literary texts become a great testing pattern to understand what NLP can and can't recognize particularly around emotions like love, grief, and ambiguity in both ancient and contemporary stories. This study focuses on whether NLP models based on the Transformer architecture can pick up on romantic undertones in the scenes depicting Patroclus' death in both versions, ultimately probing the potential and limits of NLP in interpreting literature.

Research Objectives

This study aims to:

1. Evaluate whether Transformer models can classify The Iliad's depiction of Achilles' grief as romantic or ambiguous.
2. Compare The Iliad and The Song of Achilles to quantify their emotional and linguistic differences using sentiment analysis and topic modeling.
3. Analyze the limitations of NLP in detecting nuanced and historically ambiguous relationships in literature.

Methodology

Data and Preprocessing

To facilitate a meaningful and fair comparison between the emotional portrayal of Patroclus' death in *The Iliad* and *The Song of Achilles*, this study extracted parallel narrative sections from each text. For *The Iliad*, I selected passages from Book 17, Book 18, and Book 23—the moments covering Patroclus' decision to enter battle, Achilles' grief upon learning of his death, and the final mourning rituals. This selection was based on two key criteria: first, these books contain the core canonical structure of the death event and its emotional aftermath in the original Homeric tradition; second, their narrative ambiguity, especially in Achilles' reaction, serves as fertile ground for testing the capability of NLP models to detect subtle emotional and relational nuances. Since *The Iliad* is composed in a dense poetic register and lacks explicit emotional labeling, these books preserve the epic's complex tone, which is crucial for assessing how well transformer-based models can infer grief, love, or ambiguous attachment.

For *The Song of Achilles*, I extracted text from Chapters 31, 32, and 33. These chapters represent the emotional climax of the novel and offer a modern reinterpretation of the same event from a first-person perspective, told through Patroclus' own voice. This choice reflects both narrative symmetry and thematic clarity: the author, Madeline Miller, foregrounds the romantic and psychological dimensions of their relationship, which allows a direct comparison with the more restrained and ambiguous expressions found in *The Iliad*. Additionally, Miller's contemporary language style and explicit emotional cues present an ideal contrast for testing how transformer models behave when processing classical versus modern literary emotion.

To ensure equitable model input and avoid length bias, I standardized both passages to approximately 2000 words. This equalization helps prevent the longer text from overwhelming sentiment or relationship classification outputs.

For preprocessing, I applied a minimal yet critical approach: I retained all original punctuation (especially periods, exclamations, and quotation marks), and used a regular expression to split the text into sentences while preserving emotional tone markers. I normalized whitespace to remove irregular spacing without disrupting syntactic integrity. This method supports both GoEmotions sentiment analysis and zero-shot classification by maintaining the structure and expressiveness of each sentence, which is essential for detecting fine-grained emotions like grief or love, as well as implicit relationship cues. This preprocessing preserves fidelity to each text's style while optimizing compatibility with transformer-based NLP models.

Sentiment Analysis

Introduction to transformers models and GoEmotions

Transformer-based models have really changed the way Natural Language Processing (NLP) works, making it possible for machines to understand and produce language in a way that feels more like how humans actually speak. The idea was first introduced by Vaswani et al. (2017), and what makes this model special is its use of attention mechanisms. Basically, it allows the model to figure out which words are more important in a sentence, no matter where they appear. Compared to older models like RNNs or LSTMs, Transformers can pick up on relationships between words even if they're far apart in the text. That's why they've become so popular for things like sentiment analysis, classifying texts, or even translating languages—they're just better at keeping track of the overall meaning in context.

Figure 1. Transformer's model architecture





Figure 2. The equation of attention weights

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Because of the multi-self attention (Figure 2.), transformer models are innovative and efficient tools for exploring emotion, sentiment, and character relationships within complex narratives for literary analysis. These models are trained on massive text corpora and can infer emotional cues and relational dynamics, even from unseen data (Underwood, 2019). However, whether applying transformers to literature is efficient still has potential challenges because the literary texts often contain metaphor, implicit emotion, and ambiguous expression that differ from the straightforward language typically found in social media or news data. As Bamman, Underwood, and Smith (2020) claim, while Transformer models have shown promising results in analyzing modern literary corpora, they still struggle with figurative language, subtle tone, and culturally specific narrative conventions.

Figure 3. The Google bert-based GoEmotion model architecture



For sentiment analysis, the GoEmotions model (Figure 3.) developed by Google AI, based on a BERT architecture fine-tuned on a Reddit dataset is a good tool to help me classify English text into 27 fine-grained emotion categories plus a neutral class (Demszky et al., 2020). While the model is capable of detecting clear emotional cues in everyday language, its applicability to literary texts especially ancient poetry like The Iliad still remains underexplored. Thus, in this project I adopted the GoEmotions model to quantify and visualize the emotional intensity within selected mourning scenes in order to get the data-driven perspective on the emotional trajectory of Achilles' grief.

Experiment design and sentiment analysis process

The objective of this experiment is to compare The Iliad and The Song of Achilles by analyzing the emotional intensity trajectory through their mourning scenes using sentiment analysis, which can quantify and compare the emotional and linguistic differences between two distinct literary styles. In addition, emotion intensity score analysis like neural label can also indicate whether Transformer models like GoEmotions can effectively detect emotional nuance, or whether their performance is limited when confronted with metaphorical, ambiguous, or indirect expressions of emotion common in literature.

I selected two texts for analysis in the sentiment analysis experiment:

1. The Iliad (Robert Fagles' translation), representing the original version known for its poetic ambiguity, metaphor, and restraint in emotional expression.
2. The Song of Achilles by Madeline Miller, a contemporary novel that presents a more explicit and emotionally charged portrayal of Achilles' grief and love for Patroclus.

My code is designed to preprocess each text by splitting it into manageable units (chunks) of two sentences each. The reason why I allocated the two sentences in one chunk is because the small chunk size allows the model for generating the higher emotional resolution, making it easier to capture the rise and fall of emotional intensity across the scene. The text is first cleaned using regular expressions to normalize spaces while preserving original punctuation. I then apply sentence segmentation using `re.findall()` to split the text based on sentence-ending punctuation. Once the text is divided into chunks, each chunk is processed through the GoEmotions model. The model returns a probability distribution over 28 emotion labels. I focused on five key labels for this study: grief, love, sadness, anger, and neutral. These labels were chosen because they are most relevant to the thematic content of Achilles' mourning and have been central in literary discussions regarding his relationship with Patroclus.

By storing and averaging the emotion scores from each chunk, I can compute the mean emotion intensity for each label across the entire scene. The higher the score, the more confident the model is that the corresponding emotion dominates that text segment. These selected emotions reflect both the thematic core of the mourning scenes and the reference to compare emotional trajectories across the two texts. Then I mapped the emotional trajectory by those selected emotions intensity to visualise how different emotions rise and fall across the two different literary narrative styles.

Most importantly, the neutral label intensity in (Table 1) serves as an analytical and key signal in this experiment. In Transformer models, a consistently high neutral score often indicates that the model fails to detect specific emotions due to literary complexity such as metaphor, implicit emotion, or cultural conventions of indirect speech. Therefore, if The Iliad exhibits a higher proportion of neutral scores compared to The Song of Achilles, it would support the hypothesis that Transformer-based models struggle with ancient poetic and ambiguous narrative texts due to their literary nuance.

Table 1: The emotion labels intensity in The Iliad

	grief	love	sadness	anger	neutral	chunk_index
0	0.001936	0.000947	0.011485	0.013449	0.619676	0
1	0.002634	0.000688	0.029621	0.131177	0.091987	1
2	0.006356	0.001306	0.106789	0.044532	0.069494	2
3	0.002083	0.004930	0.018417	0.003773	0.433725	3
4	0.005776	0.006560	0.030890	0.051699	0.626112	4
5	0.013787	0.002111	0.498710	0.006627	0.060267	5
6	0.000302	0.001339	0.001103	0.004402	0.908378	6
7	0.000995	0.025229	0.004234	0.011244	0.617305	7
8	0.000760	0.000825	0.004981	0.010682	0.694830	8
9	0.000174	0.000507	0.000669	0.002543	0.966775	9

Experiment objective and interpretation

The purpose of applying sentiment analysis in this study is not merely comparing The Iliad and The Song of Achilles to quantify their emotional and linguistic differences using sentiment analysis but also evaluate how well Transformer-based models perform in analyzing literature with varying degrees of emotional explicitness. By projecting the emotional trajectory through label intensities, I can get a dynamic comparison of how the two texts represent Achilles' grief. If The Song of Achilles shows higher intensity scores for love, grief, and sadness, it will confirm its explicit emotional style. Conversely, if The Iliad demonstrates a higher neutral score or lower emotional intensities across these labels, it will indicate that the Transformer model encounters difficulties in detecting specific emotions due to literary ambiguity, metaphorical language, and indirect expression.


Thus, this experiment uses sentiment analysis not only to classify emotions but also to reveal the limitations of NLP in analyzing literature. Transformer models, while powerful, operate most effectively on texts with explicit sentiment markers. Their performance diminishes in the face of poetic indirection, making The Iliad an ideal test case for evaluating the interpretative boundaries of NLP in literary analysis.

Relationship Classification and Topic modelling

Principle of Zero-shot Learning and BART-large-MNLI Model

Transformer-based language models have revolutionized NLP tasks by enabling machines to perform text classification, sentiment detection, and question answering with remarkable accuracy. One of the most important developments in recent years is Zero-shot Learning, which makes models classify texts into labels that they haven't been explicitly trained on. Instead of requiring task-specific labeled data, zero-shot learning leverages large-scale pretrained language models that generalize knowledge from their training corpora. This experiment applies the BART model fine-tuned on Facebook's large-MNLI dataset in (Figure 4.), which is a Transformer-based architecture combining not only bidirectional encoding like BERT but also autoregressive decoding like GPT. The BART model is pretrained on a denoising autoencoder task and then fine-tuned on the Multi-Genre Natural Language Inference (MNLI) dataset. The MNLI dataset consists of sentence pairs labeled as entailment, neutral, or contradiction. Using this framework, the zero-shot classification task reformulates text classification into an NLI (Natural Language Inference) problem by asking the model: Does this text entail the hypothesis (label)?

Figure 4. The BART model architecture



The core idea is that if a given input sentence strongly matches a label (like "romantic relationship"), the model will output a high probability score for the predefined label. This principle makes BART-large-MNLI particularly suitable for relationship classification in literature, where we aim to infer the underlying emotional or relational meaning within a given text passage. Therefore, because of the properties, the Zero-shot learning provides me with a powerful, data-efficient tool to analyze literature without the need for manually annotated datasets, helping us to do scalable and reproducible experiments across different texts.

Experiment objective and the design with implement process

In this experiment, I used the zero-shot classification from the Hugging Face transformers library to analyze how an NLP model interprets Achilles' grief in *The Iliad*. I applied the pre-trained facebook/bart-large-mnli model, which allows text to be classified into different relationship categories without requiring specific training on the text itself.

For this experiment, I selected five relationship labels: "romantic relationship," "brotherhood," "friendship," "devotion," and "warrior bond." These five labels represent different possible ways of understanding the emotional connection between Achilles and Patroclus. The model then compares the input text with each label and returns a probability score between 0 and 1. The possibility score means how confident a text has been allocated to defined labels. This is to say, a higher score means the model is more confident that the label matches the emotional content of the text, while a lower score means the model finds weaker evidence for that label.

After loading the mourning scene from *The Iliad*, I then inputted the text into the model with the five selected labels. And the output showcases a list of labels ranked from the highest to lowest score. For example, if the label "brotherhood" receives the highest score, it suggests that the model interprets Achilles' grief mainly as heroic companionship rather than romantic love.

To clearly present the results, I created a bar chart using matplotlib. The chart shows each label on the x-axis and its probability score on the y-axis. This visualization allows me to compare the relationship labels easily and observe how the model interprets the emotional meaning within the text based on its scoring system.

Experiment interpretation and significance

After processing the experiment, I can have a comprehensive and comparative frame for analyzing how modern NLP models interpret relationships in literature. Specifically, it enables me to evaluate the model's performance in detecting explicit vs. implicit emotional bonds across two very different literary styles.

For *The Song of Achilles*, I expect a high probability score for "romantic relationship," reflecting its modern narrative voice and explicit emotional language.

In contrast, for *The Iliad*, a higher score for "brotherhood," "devotion," or "warrior bond" (non-romantic relationship labels) along with a low score for "romantic relationship" would tend to interpret Achilles' grief through a lens of heroic friendship rather than romantic attachment, showing the model's difficulty in detecting literary nuance. Such the expected result would highlight the limitations of Transformer-based models when confronted with metaphor, indirect emotional expression, and culturally specific narrative conventions found in ancient literature. Thus, it will align with Underwood's (2019) observation that NLP models still struggle to capture figurative language and layered emotional meaning.

Thus, this experiment serves two primary purposes:

1. Evaluate whether zero-shot relationship classification can effectively differentiate between the emotional framing in *The Iliad* and *The Song of Achilles*

2. Analyse the romantic relationship score against non-romantic labels in an ambiguous text like The Iliad in order to reflect limitations of NLP models in capturing emotional subtlety, metaphor, and figurative language in literature.

Topic Modeling

Based on the zero-shot classification approach using the BART-large-MNLI transformer for relationship classification, I explored the thematic composition of Achilles' mourning scenes in The Song of Achilles by the same approach as well. The main objective was to assess the presence and intensity of predefined themes such as "love story," "grief," "honor," "destiny," "war," "loss," and "fate" without prior fine-tuning or manual annotation.

To implement this, I loaded the cleaned mourning scene text as a single input and passed it to the zero-shot-classification pipeline. This allowed the model to assess the likelihood of each candidate theme being contextually present in the passage. Compared to traditional topic modeling methods, this approach is flexible in handling long and nuanced literary narratives and can directly evaluate semantic alignment with user-defined labels.

The classification output showed a clear emphasis on "love story" and "grief," which aligns with the expected emotional depth of the passage. Specifically, "love story" received the highest score, suggesting that the model successfully picked up the emotional and romantic undertones embedded in the scene. This is particularly significant because it provides empirical evidence that the romantic bond between Achilles and Patroclus is linguistically perceivable, even without explicitly stating it.


By using this method, I was able to ignore the limitation of unsupervised topic modeling (like LDA), which may struggle with short texts or require multiple documents to identify coherent themes. In contrast, zero-shot classification provides direct thematic probabilities, allowing a straightforward comparison and interpretation between texts.

Finding and Conclusion

Sentiment Analysis


The sentiment trajectory plots and average emotion intensity plot offer critical insights for me to analyse how the Transformer-based NLP model interprets emotional tones within *The Iliad* and *The Song of Achilles*. After doing the comparative experiments, the (Figure 5. & Figure 6.) show subtle yet meaningful differences, particularly in how explicit and implicit emotional expressions influence the NLP model's interpretative capacity.

Figure 5. The average emotion intensities plot - The Iliad



In both texts, the "neutral" emotion label dominates significantly, showing average intensity scores notably higher than any specific emotional label such as sadness, love, grief, or anger. This neutral intensity is slightly higher in *The Iliad*, where the "neutral" category overwhelmingly surpasses other emotions, reflecting the NLP model's difficulty in precisely assigning emotional labels within the ancient epic's complex metaphorical language. This circumstance highlights one of the central challenges for NLP in literary studies: ancient literary texts with metaphorical and culturally nuanced expressions often fail to yield explicit emotional categorizations recognizable by a modern, literal-trained algorithm.

Figure 6. The average emotion intensities plot - The song of Achilles



Conversely, in (Figure 6.) while The Song of Achilles is still identified as "neutral" as the strongest emotional category, showing slightly elevated scores for sadness, love, and anger compared to The Iliad. This subtle increase suggests that the NLP model more effectively captures explicit emotional statements typical of contemporary literary styles.

Figure 7. The emotion trajectory plot - The Iliad





Figure 8. The emotion trajectory plot - The song of Achilles




The trajectory plots in (Figure 7. & Figure 8.) further clarify these differences. In *The Iliad* (Figure 7), emotional intensities aside from neutral tend to stay low and don't show any major spikes. This ongoing neutrality suggests that the model struggles to confidently pick up on Homer's metaphorically and emotionally layered expressions. On the other hand, in *The Song of Achilles* (Figure 8.), some emotional categories like sadness and love display sharper and more significant peaks, indicating clear narrative instances where the author explicitly describes emotions. For instance, the "love" label in Miller's text shows occasional distinct intensity peaks due to the explicit modern emotional language. Although these peaks remain modest relative to the neutral baseline, they still indicate the NLP model finds it is relatively easier to recognize clear emotion in contemporary language.

However, even in Miller's more explicit emotional scenes, the prevalence of high numbers of neutral classifications still pose a fundamental issue. Although the clear style of modern narrative helps NLP to recognise emotions more easily, it does not eliminate and solve difficulties of literary nuance, metaphor, and implied emotions. Even though Miller's contemporary adaptation is much clearer in style, the model still assigns a lot of "neutral" labels, which shows that it often struggles to catch the more delicate emotional layers. Based on these results, it's clear that Miller's modern version gives the model stronger emotional cues to work with, while Homer's original text filled with metaphor and subtle shifts in tone is much harder for the model to interpret. The high frequency of neutral classifications, especially in *The Iliad*, highlights a core limitation in how NLP handles literary language. Because when emotion is implied rather than directly stated, the model tends to miss it. This sentiment analysis shows that clarity in writing plays a big role in how accurately the model can detect emotional tone, reminding us that we need to be cautious when applying NLP tools to literature, especially with texts that rely heavily on poetic expression and nuance.

NLP detection of romantic undertones and implications


In order to distinguish the relationship classification from two texts, I leveraged the zero-shot classification and topic modeling analyses to illustrate how effectively NLP models can detect romantic undertones in literary texts and how this ability varies significantly between ancient and modern narratives.

Figure 9. The zero- classification- The song of Achilles



The results from (Figure 9. & Figure 10.) clearly distinguish between the relational dynamics depicted in *The Iliad* and *The Song of Achilles*. For *The Song of Achilles*, the NLP model assigns a notably high probability score (approximately 0.90) to the label "romantic relationship," reflecting Miller's explicit narrative framing of Achilles and Patroclus' bond as romantic. Other relational labels such as "warrior bond," "brotherhood," and "devotion" also receive high scores, highlighting the complex, multi-dimensional portrayal of their relationship. Thus, the high probability score of the "romantic relationship" category validates the NLP's ability to recognize clear modern textual signals of romantic love.

Figure 10. The zero-classification plot - Iliad



By contrast, in *The Iliad*, the higher probability scores appear in categories like "warrior bond," "devotion," and "brotherhood," while the "romantic relationship" label has a notably lower score (around 0.40). The strong probability of non-romantic labels tend to align with the epic's historical and literary tradition, reflecting Homer's depiction of their relationship primarily through mutual loyalty, fraternity, and destiny rather than explicit romantic language. The model's classification in this case inclines to emotional ambiguity and confirms that when the literature doesn't have overt textual markers, NLP tends to classify the relationship as cultural frames such as warrior bond, brotherhood and devotion in the *Iliad*.

The topic modeling results (Figure 11 and Figure 12) further support this finding. In *The Iliad*, the main themes that stand out are loss, grief, war, fate, and destiny—fitting the traditional tragic and heroic tone of the epic. The low presence of "love story" aligns with what was seen in the zero-shot classification, confirming that the text focuses more on honor and fatalism than romantic emotion. In contrast, *The Song of Achilles* brings in similar themes like fate and destiny, but "love story" appears much more prominently. After implementing the topic modelling, the result shows that the contemporary literature where the romantic love is explicitly expressed will reinforce the NLP to classify the relationship.

Figure 11. The Topic Modelling - The Iliad





Figure 12. The Topic Modelling - The song of Achilles



Bringing both methods together, the outcome showcases that Transformer-based models are capable of detecting romantic relationships only when the narrative is explicit and direct such as Miller's adaptation. However, in *The Iliad*, where emotion is layered in metaphor and literary nuance rather than clearly stated, the models will fail to identify Achilles' grief as a romantic relationship because the romantic score is low, and the model regards the grief more through themes of heroism or brotherhood. Based on this outcome, the experiment shows that Transformer models have difficulty identifying the grief in *The Iliad* as romantic, regarding it as ambiguous and interpreting it in non-romantic relational categories like loyalty or brotherhood. Thus, based on the research findings, I believe that NLP still relies heavily on clear language and familiar cultural patterns to interpret emotional nuance. When love is expressed through poetic or symbolic forms like in *The Iliad*, the Transformer model will struggle to detect it, often failing to recognize the romantic elements altogether.

NLP in humanities analysis : strengths and limitations

Comparing the two texts side by side has helped me better understand where NLP offers real value in literary research and where its limitations become more apparent. One clear strength is when we apply zero-shot classification and topic modeling, which model can efficiently identify emotional tones or relationship cues when the literature is contemporary and explicit. In *The Song of Achilles*, for instance, the model was able to detect romantic undertones and emotional shifts with a high degree of accuracy. This demonstrates that NLP can serve as a useful tool for digital humanities scholars working with contemporary literature, especially when the aim is to identify dominant emotional patterns or relationship dynamics across longer texts. When emotions are clearly stated, the model performs well and can offer a systematic, reproducible layer of interpretation that complements traditional close reading.

However, these strengths did not perform well when I applied the model to the *Iliad*. Most of the emotional content was labeled as "neutral," and the probability assigned to "romantic relationship" was noticeably lower. This suggests that the model struggled to process the more layered and symbolic expressions found in Homer's language. Despite the presence of grief and emotional intensity in the text, the model did not interpret it through a romantic lens. This isn't entirely surprising, as Homer rarely presents emotion directly. Instead, much of the emotional meaning in ancient literature is conveyed through cultural or heroic codes, which aren't easily accessible to models trained on contemporary datasets. These results also align with Underwood's (2019) observation that NLP often fails to deal with metaphor, subtle emotional cues, and culturally embedded narrative structures.

What became increasingly clear through this process is that texts like the *Iliad* express layers of emotions in ways that are difficult for NLP to detect. The meaning often depends on context, historical convention, or even what remains unsaid such as subtleties that Transformer models cannot grasp unless explicitly trained to do so. While NLP performs well with clear, literal language, it still can not perfectly work when interpreting poetic or symbolic texts. For this reason, I believe that we need to analyse the results from these transformer models with a certain level of caution because the classifications may be informative, which can be improved by fine-tuning with the annotated other ancient works

and applying prompt engineering to better capture implied emotional nuance. Overall, I think integrating NLP tools with traditional humanistic analysis is still the more balanced and insightful approach for studying texts where emotional meaning is implied rather than directly stated.

Bibliography

- Bamman, D., Underwood, T., & Smith, N. A. (2020). A Digital Humanities Approach to Emotional Arcs in Literature. Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics.
<<https://aclanthology.org/2020.acl-main.463/>>
- Cantarella, E. (1992). *Bisexuality in the Ancient World* (C. Ó. Love, Trans.). Yale University Press.
- Fagles, R. (Trans.). (1998). *The Iliad*. Penguin Books.
- Halperin, D. M. (1990). *One Hundred Years of Homosexuality: And Other Essays on Greek Love*. Routledge.
- Miller, M. (2011). *The Song of Achilles*. Ecco Press.
- Nagy, G. (1979). *The Best of the Achaeans: Concepts of the Hero in Archaic Greek Poetry*. Johns Hopkins University Press.
- Roisman, H. M. (2012). Review of *The Song of Achilles*. Bryn Mawr Classical Review. <<http://bmcr.brynmawr.edu/2012/2012-07-52.html>>
- Underwood, T. (2019). *Distant Horizons: Digital Evidence and Literary Change*. University of Chicago Press.
<<https://press.uchicago.edu/ucp/books/book/chicago/D/bo42781985.html>>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30.
<https://papers.nips.cc/paper_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf>